

AI writes it
in minutes.
I don't ship it
for *hours*.

Every non-trivial change I build passes through three independent reviewers before it ships. Always in the same order.

Speed is cheap. *Trust* is not.


AI can generate a feature in minutes. But speed with no guardrails ships bugs, security holes, and ungoverned AI straight into systems that hold real patient data.

So AI-written code never goes straight to production. **It earns its way there.**

3


Three gates.
None of them *me*.

Every change runs through three independent, zero-context reviewers, each with one job. It ships only when all three say yes.



THE PIPELINE

It's a *loop*,
not a line.



Any rejection sends the change **back to the code**, then back through the gates from the top. Nothing skips ahead.

Is the code *right?*

- Correctness, readability, performance, security on the actual diff
- Unbiased. It never wrote the code it is reviewing
- NEEDS CHANGES means I fix it and run the gate again

PASS

PASS WITH NOTES

NEEDS CHANGES

Does it *work?*

- Generates and runs tests: happy path, edge cases, errors
- A clean PASS is the only way through
- On FAIL I fix the code, not the test

PASS

FAIL

RUNS ONLY AFTER CODE-REVIEWER PASSES

Is it *safe* and accountable?

The last gate, and the one I never let a build skip.

- Scored against an 8-layer AI governance framework I adapted and operate
- Enforcement weighted by the product's risk tier
- Catches shadow AI, AI on sensitive data, decisions with no human sign-off

RUNS ONLY AFTER CODE AND QA PASS

Right. Then working. Then *governed*.

- 1 Code first. No point testing logic that is wrong by design.
- 2 QA second. Prove the reviewed code actually behaves.
- 3 Governance last. You can only govern a thing that works.

Each gate assumes the one before it passed. Reorder them and you waste the expensive checks.

The difference between a demo and a *system*.

Real products, not toys. This loop is why the things I ship hold up where it counts:

- Production-grade code, not code that ran once
- Real test coverage on every change
- AI that is governed before it ever touches patient data

BUILD WITH ME

This is how I *ship*.

Not the fastest way to deploy. The way I can stand behind, on every build.

I share the methods behind everything I build, **in public**. Follow along.

Building something that has to be trusted? **I'd be glad to talk.**